

INDEX

S.No	Title	Page No.	Date	Remarks
1	Construct various type of plots/charts like histogram, bar chart, pie-chart and scatter plot by importing data from a CSV format file. Further label different axes and data in a plot.	2	21/07/2023	
2	Perform data cleaning and data preprocessing tasks.	4	28/07/2023	
3	Perform linear regression on boston datasets.	5	04/08/2023	
4	Perform multiple linear regression on boston datasets.	6	11/08/2023	
5	Perform logistics regression on generated datasets	7	18/08/2023	
6	Write a program to train the Naïve Bayes classifier for a sample training data set stored as a .csv file. Compute the accuracy of the classifier, considering the test dataset.	9	25/08/2023	
7	Write a program to implement K-Nearest Neighbour algorithm to classify the sample data set into different classes. Print both correct and wrong predictions.	10	01/09/2023	
8	Write a program to train a Decision Tree Classifier to classify the sample dataset into different categories.	11	15/09/2023	
9	Implement K- means clustering on wine datasets.	12	06/10/2023	
10	Write a program to implement Principal Component Analysis to reduce the number of dimensions and thereby the size of the raw dataset.	13	13/10/2023	

LAB 1

Aim: Construct various type of plots/charts like histogram, bar chart, pie-chart and scatter plot by importing data from a CSV format file. Further label different axes and data in a plot.




Program:

```
import random
import pandas as pd
import matplotlib.pyplot as plt
data={'age':[random.randint(20,60) for _ in
range(100)],'gender':[random.choice(["Male","Female"]) for _ in
range(100)],'income':[random.randint(20000,100000) for _ in range(100)]}
df=pd.DataFrame(data)
plt.hist(df['age'])
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Age Distribution')
plt.show()
plt.bar(df['gender'].unique(), df['gender'].value_counts())
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Gender Comparison')
plt.show()
print(type(df['gender'].value_counts()))
print(df['gender'].value_counts())
print(type(df['gender'].unique()))
print(df['gender'].unique())
import matplotlib.pyplot as plt
labels = 'Frogs', 'Hogs', 'Dogs', 'Logs'
sizes = [15, 30, 45, 10]



fig, ax = plt.subplots()
ax.pie(sizes, labels=labels)
plt.pie(df['gender'].value_counts(), labels=df['gender'].unique())
plt.title('Gender Proportion')
plt.show()

plt.scatter(df['age'], df['income'])
plt.xlabel('Age')
plt.ylabel('Income')
plt.title('Age vs Income')
plt.show()
```

Output:



Gender Proportion



LAB 2

Aim: Perform data cleaning and data preprocessing tasks.

Program:

```
import pandas as pd
import numpy as np
df = pd.read_csv('/content/data.csv')
print(df.columns)
print(df.isnull().sum())
df['income'].fillna(df['income'].mean(), inplace=True)
df['date'].fillna(method='ffill', inplace=True)
df.dropna(inplace=True)
df = df.drop('gender', axis=1) # Remove irrelevant column
df = df.assign(age_squared=lambda x: x['age']**2) # Insert new column with
age squared
df = df.rename(columns={'income': 'annual_income'}) # Rename target variable
column
print(df.columns)
from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler
scaler = MinMaxScaler()
df['age'] = scaler.fit_transform(df[['age']])
scaler = StandardScaler()
df[['annual_income', 'age_squared']] =
scaler.fit_transform(df[['annual_income', 'age_squared']])
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['marital_status'] = le.fit_transform(df['marital_status'])
print(df)
```

Output:

```
Index(['age', 'income', 'date', 'marital_status', 'gender'], dtype='object')
age          0
income       1
date         1
marital_status  0
gender       0
dtype: int64
```

```
Index(['age', 'annual_income', 'date', 'marital_status', 'age_squared'],
dtype='object')
```

	age	annual_income	date	marital_status	age_squared
0	0.545455	-0.313112	2020-01-01	0	-0.128298
1	0.818182	0.939336	2020-01-02	1	0.827650
2	0.000000	-1.565561	2020-01-03	0	-1.587378
3	0.409091	-0.939336	2020-01-04	1	-0.549671
4	0.227273	0.000000	2020-01-05	0	-1.052802
5	0.590909	0.000000	2020-01-06	1	0.020545
6	1.000000	1.565561	2020-01-07	0	1.548805
7	0.727273	0.313112	2020-01-07	1	0.492230
8	0.954545	1.252449	2020-01-09	0	1.362227
9	0.272727	-1.252449	2020-01-10	1	-0.933308


LAB 3

Aim: Perform linear regression on boston datasets.

Program:

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
df=pd.read_csv("/content/drive/MyDrive/Class/My_lab/Lab4/Boston.csv")
X=df[['rm']]
y=df['medv']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
random_state=42)
model=LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
import matplotlib.pyplot as plt
plt.scatter(X_train, y_train)
plt.plot(X_test,y_pred,color='red')
plt.xlabel('Actual target values')
plt.ylabel('Predicted target values')
plt.show()
```

Output:



LAB 4


Aim: Perform multiple linear regression on boston datasets.

Program:

```
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
df=pd.read_csv("/content/drive/MyDrive/Class/My_lab/Lab5/Boston.csv")
X=df[['crim','zn','indus','chas','nox','rm','age','dis','rad','tax','ptratio',
      'black','lstat']]
y=df['medv']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5,
random_state=42)
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
import matplotlib.pyplot as plt
plt.scatter(y_test,y_pred)
lims = [min(min(y_test), min(y_pred)), max(max(y_test), max(y_pred))]
plt.plot(lims, lims, 'k--')
from sklearn.metrics import r2_score
print(r2_score(y_pred,y_test))
```

Output:

0.6253969844551934



LAB 5

Aim: Perform logistics regression on generated datasets.

Program:

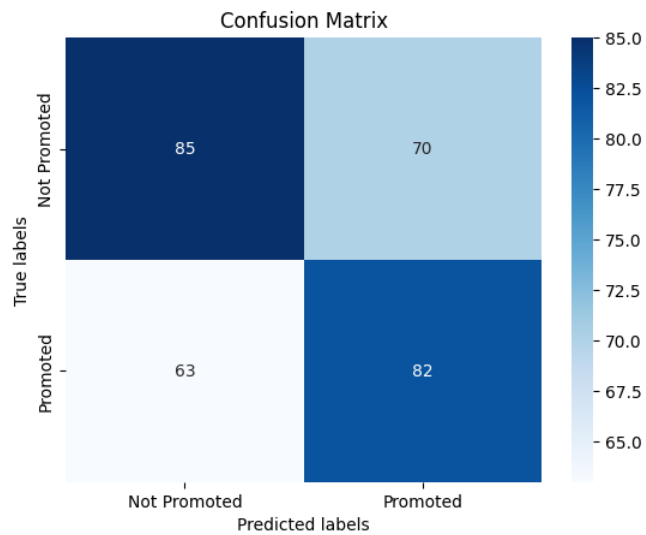
```
import numpy as np
import pandas as pd
age=np.random.normal(40,10,n)
gender=np.random.choice(["Male","Female"],n)
education = np.random.choice(['high school', 'college', 'graduate'], n)
job_level = np.random.choice(['junior', 'senior'], n)
last_evaluation = np.random.uniform(0.4, 1, n)
average_monthly_hours = np.random.randint(100, 300, n)
time_spend_company = np.random.randint(1, 10, n)
number_of_projects = np.random.randint(1, 7, n)
work_accident = np.random.choice([0, 1], n)
promotion = np.random.choice([0, 1], n)
salary = np.random.choice(['low', 'medium', 'high'], n)
dictionary={'age':age,'gender':gender,'education':education,'job_level':job_level,
'last_evaluation':last_evaluation,'average_monthly_hours':average_monthly_hours,
'time_spend_company':time_spend_company,'number_of_projects':number_of_projects,
'work_accident':work_accident,'promotion':promotion,'salary':salary}
df=pd.DataFrame(dictionary)
data=pd.get_dummies(df,columns=['gender','education','job_level','salary'])
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix,classification_report
from sklearn.model_selection import train_test_split
model=LogisticRegression()
X=data.drop('promotion',axis=1)
Y=data['promotion']
train_x,test_x,train_y,test_y=train_test_split(X,Y,test_size=0.3)
model.fit(train_x,train_y)
y_pred=model.predict(test_x)
print(classification_report(test_y, y_pred))
print(confusion_matrix(test_y, y_pred))
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(test_y, y_pred)
ax = plt.subplot()
sns.heatmap(cm, annot=True, ax=ax, cmap='Blues', fmt='g')
ax.set_xlabel('Predicted labels')
ax.set_ylabel('True labels')
ax.set_title('Confusion Matrix')
ax.xaxis.set_ticklabels(['Not Promoted', 'Promoted'])
ax.yaxis.set_ticklabels(['Not Promoted', 'Promoted'])
plt.show()
```

Output:

```
precision    recall  f1-score   support

           0       0.57    0.55    0.56    155
           1       0.54    0.57    0.55    145

 accuracy          0.56    300
 macro avg          0.56    300
 weighted avg       0.56    300
[[85 70]
 [63 82]]
```



LAB 6

Aim: Write a program to train the Naïve Bayes classifier for a sample training data set stored as a .csv file. Compute the accuracy of the classifier, considering the test dataset.

Program:

```
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix
spam_df = pd.read_csv('/content/drive/MyDrive/Class/My_lab/Lab 7 Naive
Bayes/spam.csv', encoding='latin-1')
spam_df = spam_df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1)
spam_df = spam_df.rename(columns={'v1': 'label', 'v2': 'text'})
train_data = spam_df[:4400]
test_data = spam_df[4400:]
vectorizer = CountVectorizer()
train_vectors = vectorizer.fit_transform(train_data["text"])
test_vectors = vectorizer.transform(test_data["text"])
nb_classifier = MultinomialNB()
nb_classifier.fit(train_vectors, train_data["label"])
predictions = nb_classifier.predict(test_vectors)
accuracy = accuracy_score(test_data["label"], predictions)
print("Accuracy:", accuracy)
print("Confusion Matrix:")
print(confusion_matrix(test_data["label"], predictions))
```

Output:

```
Accuracy: 0.9863481228668942
Confusion Matrix:
[[1015    8]
 [   8  141]]
```

LAB 7

Aim: Write a program to implement K-Nearest Neighbour algorithm to classify the sample data set into different classes. Print both correct and wrong predictions.

Program:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import r2_score
data=load_iris()
x_train,x_test,y_train,y_test=train_test_split(data.data,data.target,test_size=0.2)
model=KNeighborsClassifier(n_neighbors=3)
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
print(r2_score(y_test,y_pred))
for i in range(len(y_pred)):
    if y_pred[i] != y_test[i]:
        print(f"Wrong prediction: Actual label = {y_test[i]}, Predicted label = {y_pred[i]}")
```

Output:

```
0.8861138861138861
Wrong prediction: Actual label = 1, Predicted label = 2
Wrong prediction: Actual label = 2, Predicted label = 1
Wrong prediction: Actual label = 2, Predicted label = 1
```

LAB 8

Aim: Write a program to train a Decision Tree Classifier to classify the sample dataset into different categories.

Program:

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X=load_iris().data
y=load_iris().target
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=2)
dt=DecisionTreeClassifier()
dt.fit(x_train,y_train)
y_pred=dt.predict(x_test)
print("Accuracy: ",accuracy_score(y_test,y_pred))
from sklearn.model_selection import GridSearchCV

params = {'max_depth': [1, 2, 3, 4, 5]}
clf = DecisionTreeClassifier()
grid_search = GridSearchCV(clf, param_grid=params, cv=5)
grid_search.fit(x_train, y_train)
print("Best parameters:", grid_search.best_params_)
```

Output:

Accuracy: 0.9333333333333333

Best parameters: {'max_depth': 5}

LAB 9

Aim: Write a program to implement K-means algorithm to form clusters in a sample of datasets on wine dataset.

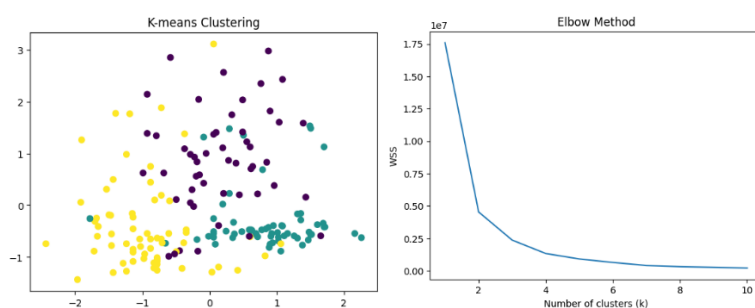
Dataset description: It contains 178 observations of wine grown in the same region in Italy. Each observation is from one of three cultivars (the 'Class' feature), with 13 constituent features that are the result of a chemical analysis.

Program:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
df=pd.read_csv("/content/wine.data",header=None)
df=df.drop(0,axis=1)
print(df.shape)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
wine_data_scaled = scaler.fit_transform(df)
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(wine_data_scaled)
plt.scatter(wine_data_scaled[:, 0], wine_data_scaled[:, 1],
c=kmeans.labels_)
plt.title('K-means Clustering')
plt.show()
wss_values = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(df)
    wss_values.append(kmeans.inertia_)
# Plot the WSS values against different values of k
plt.plot(range(1, 11), wss_values)
plt.title('Elbow Method')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WSS')
plt.show()
```

Output:

(178, 13)



LAB 10

Aim: Write a program to implement Principal Component Analysis to reduce the number of dimensions and thereby the size of the raw dataset.

Program:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
# set seed for reproducibility
np.random.seed(42)
data = np.random.randn(1000, 10) # generate random data
# create DataFrame with column names
df = pd.DataFrame(data, columns=['feature_1', 'feature_2', 'feature_3',
'feature_4', 'feature_5', 'feature_6', 'feature_7', 'feature_8',
'feature_9', 'feature_10'])
df['label'] = np.random.randint(0, 2, size=1000) # add labels to DataFrame

print(df.shape)
X = df.drop('label', axis=1)
y = df['label']
scaler = StandardScaler()
X = scaler.fit_transform(X)
pca = PCA()
X_pca = pca.fit_transform(X)
explained_variance = pca.explained_variance_ratio_ #Determine the number of
principal components to retain
n_components = np.argmax(np.cumsum(explained_variance) >= 0.90) + 1
pca = PCA(n_components=n_components) #Transform the features using the
retained principal components
X_pca = pca.fit_transform(X)
new_df = pd.concat([pd.DataFrame(X_pca), y], axis=1)
new_df.to_csv('new_data.csv', index=False)
print(new_df.shape)
```

Output:

```
(1000, 11)
```

```
(1000, 10)
```